

Agent Skills

Prototype: Table Structure Recognition

Agile SoDA

DOCUMENT AI 알고리즘 팀

2026.01.19

Goal

Build Modular Agent Skills for a TSR Prototype

- Agent → Autonomous, minimizing human intervention
- Skills → Task specific modular
- TSR (Task) → Table Structure Recognition
- Prototype → Simple, end-to-end functional, modular, and extensible

Content

- **Introduction**
- **Proposed System**
- **Implementation**
- **Result**
- **Conclusion**

Introduction

■ What is Agent Skills?

- Task specific modular
- Activates only when needed
- E.g., TSR, classification, and summarization
- Dev simplified: "Folder"

Agent Skills

Agent Skills are modular capabilities that extend Claude's functionality. Each Skill packages instructions, metadata, and optional resources (scripts, templates) that Claude uses automatically when relevant.

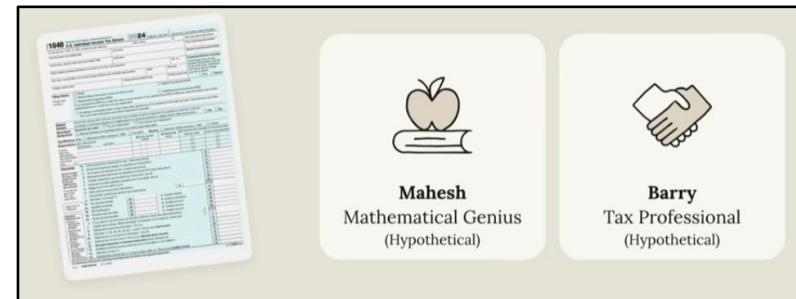
[Claude Agent Skill](#)

Skill Anatomy

```
skill-name/
├── SKILL.md
├── scripts/
└── references/
└── assets/
```

■ Why we need Agent Skills?

- Domain/Task expertise
- Modular & Extensible



[Claude Agent Skill](#)

Agent Skills

```
skills/
├── skill_1/
├── skill_2/
└── skill_3
└── ...
```

Introduction

■ Available Agent Skills

- [Claude - Agent Skills](#)
- [Google Antigravity – Agent Skills](#)
- [Cursor – Agent Skills](#)
- [OpenAI – CodeX – ExecPlan](#)
- [Google Gemini – Gems](#)

Skill Anatomy

```
skill-name/
└── SKILL.md (required)
    ├── YAML frontmatter metadata (required) (Triggering & routing - When to use)
    │   └── name: (required)
    │   └── description: (required)
    └── MD instructions (required) (Instructional control - workflow, heuristic)
        └── Bundled Resources (optional)
            ├── scripts/ - Executable code (Executable implementation - Python/Bash)
            ├── references/ - Doc loaded into context as needed (Contextual knowledge)
            └── assets/ - Files used in output (Output artifacts)
```

[Claude Agent Skill](#)

Each skill should be a folder containing a `SKILL.md` file:

```
1 .cursor/
2   └── skills/
3     └── my-skill/
4       └── SKILL.md
```

`SKILL.md` file format

Each skill is defined in a `SKILL.md` file with YAML frontmatter:

```
1 ---
2   name: my-skill
3   description: Short description of what this skill does and when to use it.
4 ---
5
6 # My Skill
7
8 Detailed instructions for the agent.
9
10 ## When to Use
11
12 - Use this skill when...
13 - This skill is helpful for...
14
15 ## Instructions
16
17 - Step-by-step guidance for the agent
18 - Domain-specific conventions
19 - Best practices and patterns
```

Every skill needs a `SKILL.md` file with YAML frontmatter at the top:

```
---
name: my-skill
description: Helps with a specific task. Use when you need to do X or Y.
---

# My Skill

Detailed instructions for the agent go here.

## When to use this skill

- Use this when...
- This is helpful for...

## How to use it

Step-by-step guidance, conventions, and patterns the agent should follow.
```

[Cursor – Agent Skills](#)

```
# Sample AGENTS.md file

## Dev environment tips
- Use `pnpm dix turbo run where <project_name>` to jump to a package instead of scanning with `ls`.
- Run `pnpm install --filter <project_name>` to add the package to your workspace so Vite, ESLint, and TypeScript can see it.
- Use `pnpm create vite@latest <project_name> --template react-ts` to spin up a new React + Vite package with TypeScript checks ready.
- Check the name field inside each package's package.json to confirm the right name—skip the top-level one.

## Testing instructions
- Find the CI plan in the .github/workflows folder.
- Run `pnpm turbo run test --filter <project_name>` to run every check defined for that package.
- From the package root you can just call `pnpm test`. The commit should pass all tests before you merge.
- To focus on one step, add the Vitest pattern: `pnpm vitest run -t <test name>`.
- Fix any test or type errors until the whole suite is green.
- After moving files or changing imports, run `pnpm lint --filter <project_name>` to be sure ESLint and TypeScript rules still pass.
- Add or update tests for the code you change, even if nobody asked.

## PR instructions
- Title format: [<project_name>] <Title>
- Always run `pnpm lint` and `pnpm test` before committing.
```

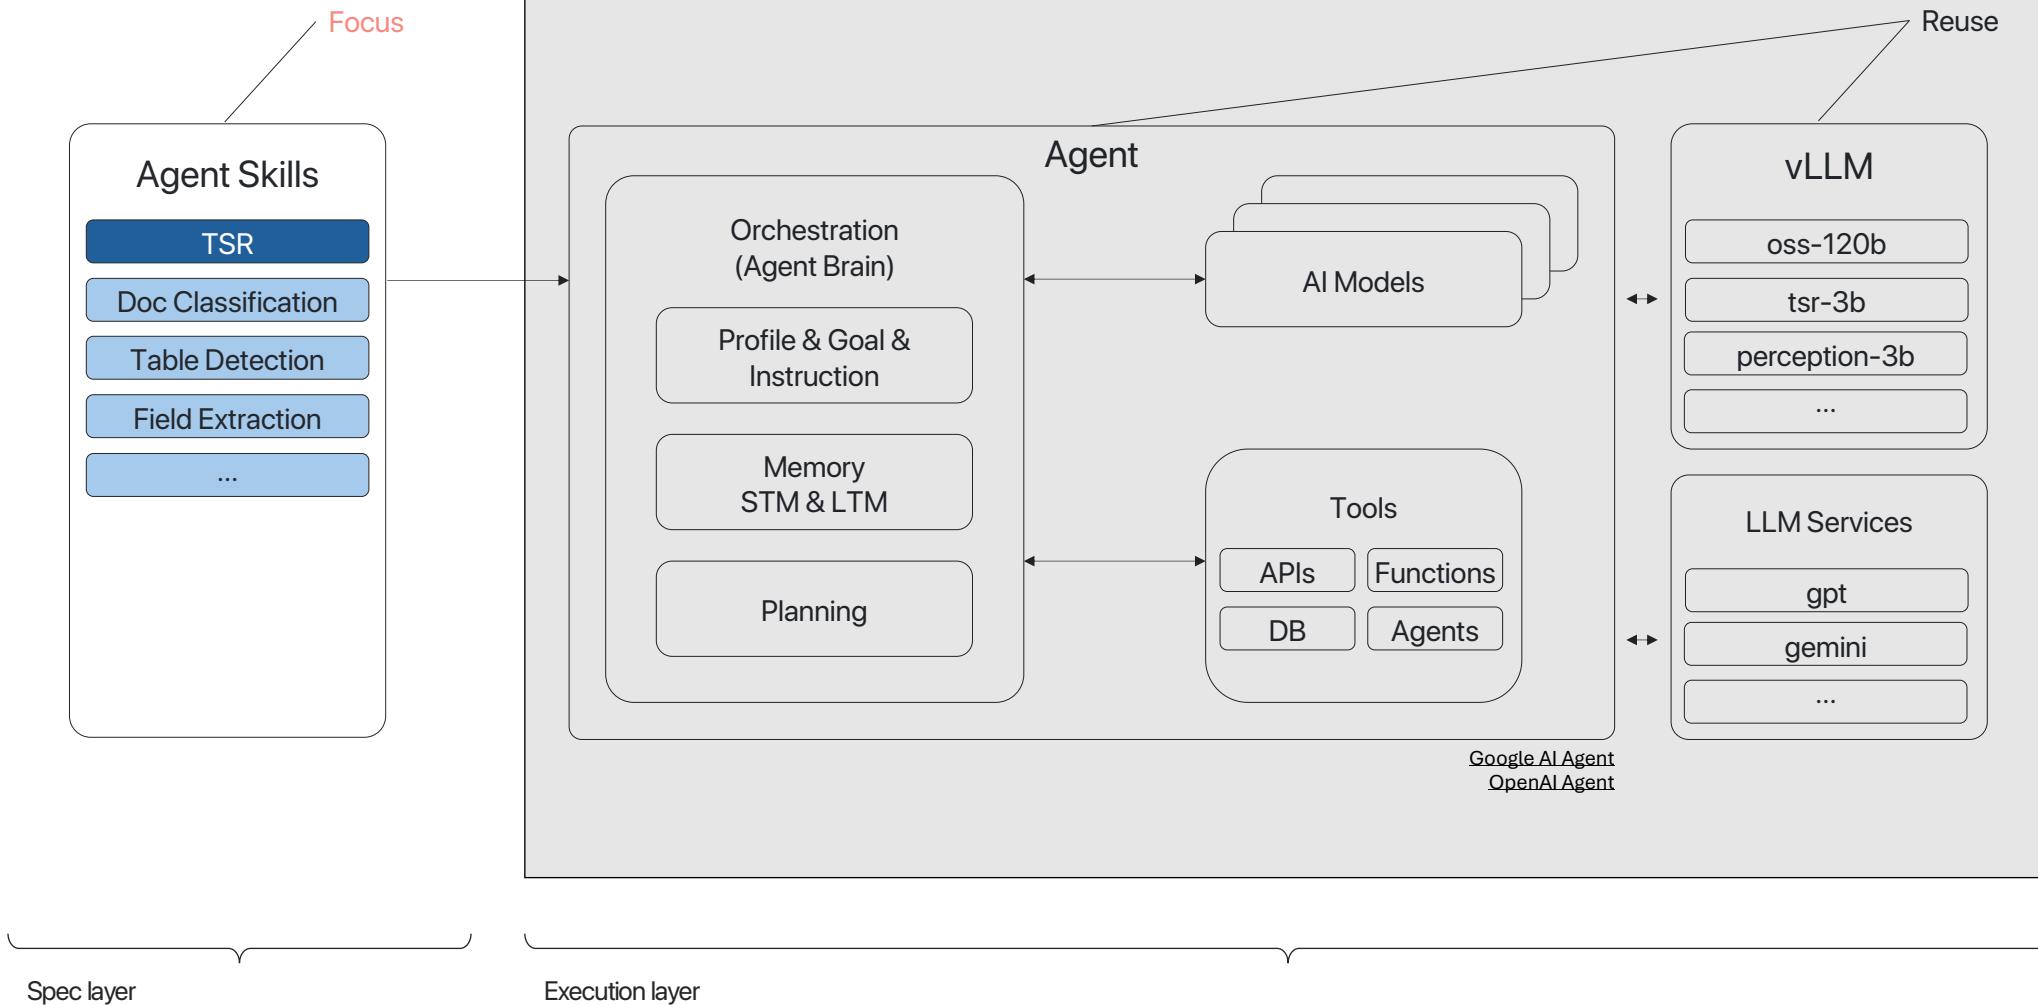
[Google Antigravity – Agent Skills](#)

[OpenAI – CodeX – ExecPlan: AGENTS.md | PLAN.md](#)

Proposed System

■ Agent Skills

- TSR
- ...



Implementation

Re-use:

- **vLLM**

- (OpenAI) oss-120b: LLM, Agent Brain
- (AgileSoDA) tsr-3b: VLM, TSR as tool (note: Finetuned VLM TSR lost Reasoning capability)
- (Nanonets) perception-3b: VLM, perception/ocr as tool

- **Agent**

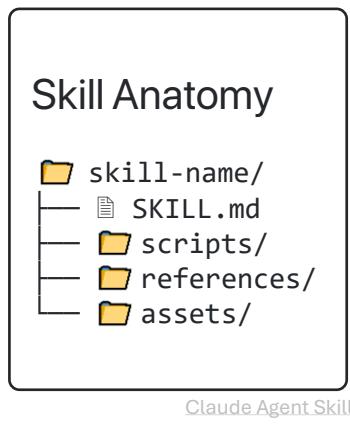
- (OpenAI) Agent (For quick prototyping)
- (AgileSoDA) Agentic Framework (Daniel, from the scratch)

Implementation

■ Agent Skills

▪ (Claude) Agent Skill

- Reuse: Concepts and Skill Anatomy
- Implementation is private



Progressive Disclosure Design Principle

How Skills work

Skills are **model-invoked**: Claude decides which Skills to use based on your request. You don't need to explicitly call a Skill. Claude automatically applies relevant Skills when your request matches their description.

When you send a request, Claude follows these steps to find and use relevant Skills:

1 Discovery

At startup, Claude loads only the name and description of each available Skill. This keeps startup fast while giving Claude enough context to know when each Skill might be relevant.

2 Activation

When your request matches a Skill's description, Claude asks to use the Skill. You'll see a confirmation prompt before the full `SKILL.md` is loaded into context. Since Claude reads these descriptions to find relevant Skills, **write descriptions** that include keywords users would naturally say.

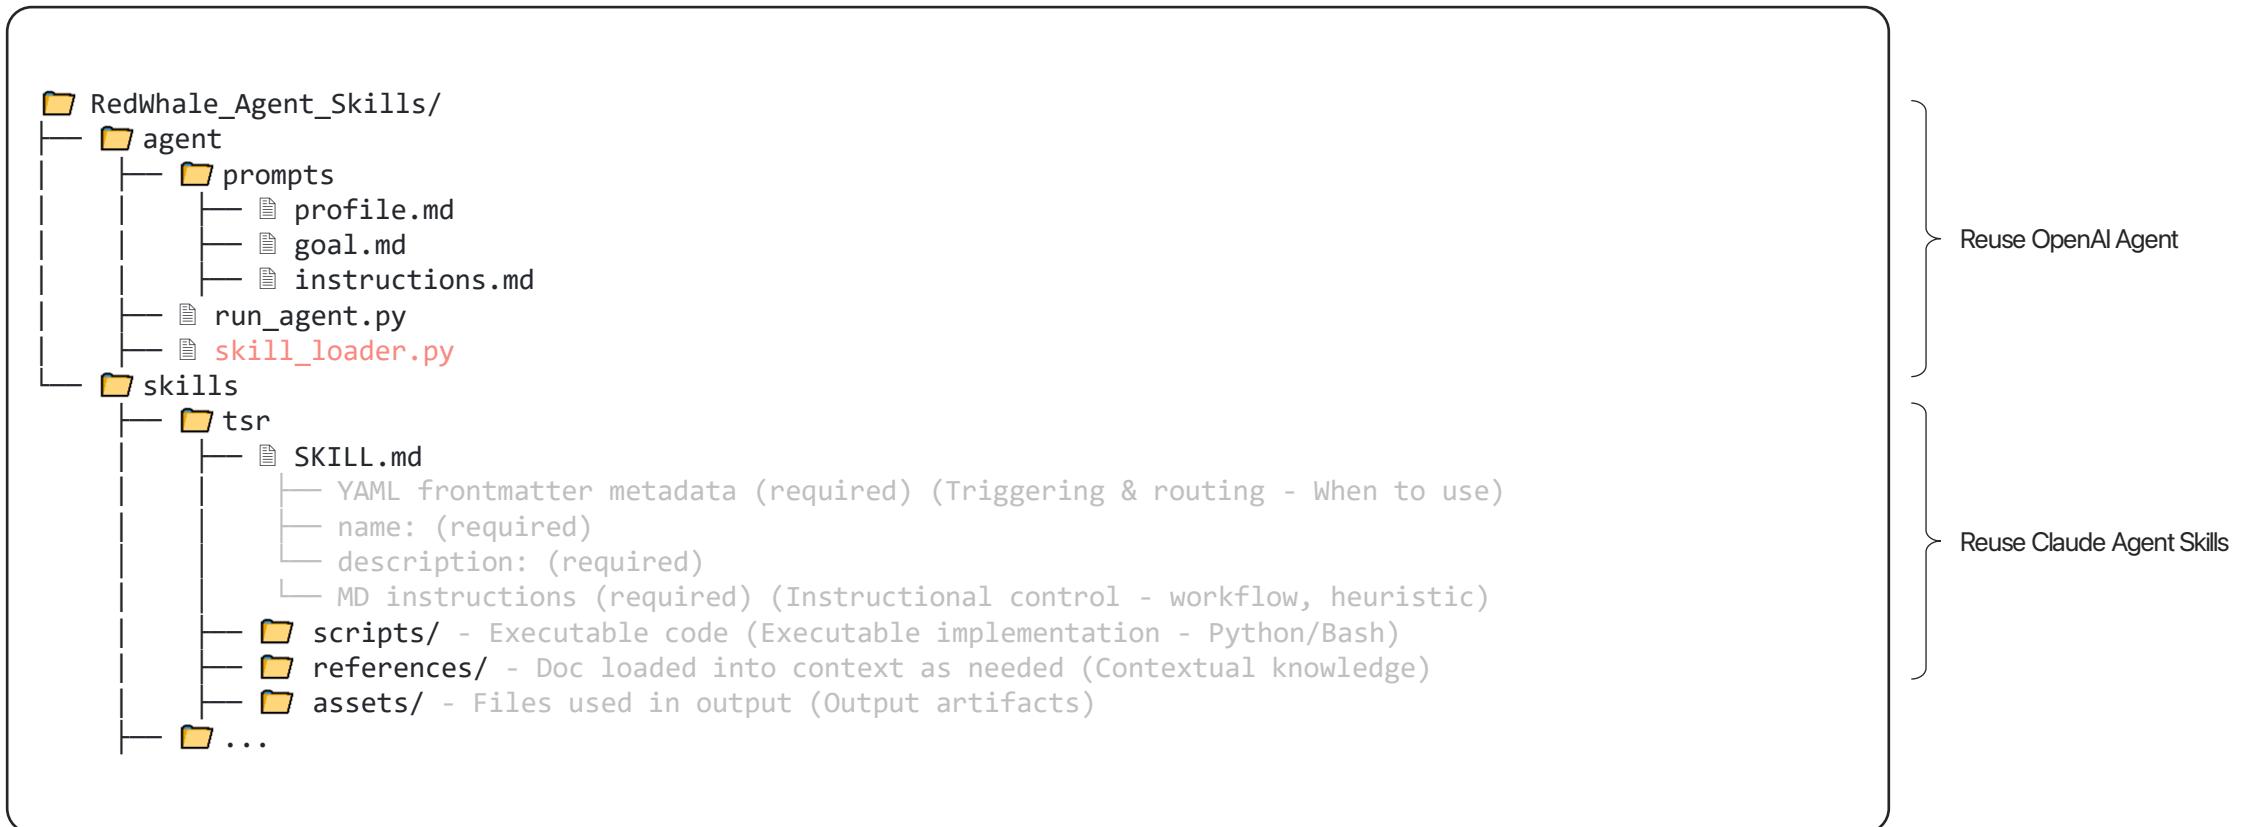
3 Execution

Claude follows the Skill's instructions, loading referenced files or running bundled scripts as needed.

- 1. Which skill is needed
- 2. Load the required skill
- 3. Execute

Implementation

■ Project Structure



Implementation

■ TSR Skill

SKILL.md

```
---
```

```
name: tsr
description: Extract structured tables from table images and return them as HTML and/or OTSL using TSR or Perception VLM backends.
---
```

```
# TSR Skill
```

```
## Workflow
```

```
### 1) Extract table from an image
- Call: `tsr_extract(image_path, backend, output_format)`
- Backends:
  - **tsr** - dedicated TSR model (best for structure accuracy)
  - **perception** - general VLM (may return HTML directly)
```

```
### 2) Detect output format
- If result contains OTSL tags → treat as **OTSL**
- Otherwise → treat as **HTML**
```

```
### 3) Convert when needed
- OTSL → HTML using `otsl_to_html(otsl_text)`
```

```
### 4) Save result (optional)
- Save HTML file with `save_html(html, path)`
```

```
## Supported Outputs
- `html` - final structured table
- `otsl` - intermediate structured format
- `auto` - return both (recommended)
```

```
## Output Contract
- **tsr_extract** → raw text + detected format + OTSL/HTML
- **otsl_to_html** → HTML
- **save_html** → file path
```

```
## Prompting Rules
- Deterministic mode (`temperature=0`)
- High `max_tokens` for large tables
- Use TSR-focused prompts
```

```
## Debugging & Recovery
- Prefer **backend="tsr"** if quality is poor
- Increase max tokens if truncated
- Re-convert from OTSL when HTML is malformed
- Ensure input image is a proper table crop
```

```
## Notes
- Use tool calls for reliability
- Keep detailed prompts in reference files
```

scripts/tsr_tools.py

```
# Exposed TSR Tools - Interfaces Only
```

```
def tsr_extract(
    image_path: str,
    backend: str = "tsr",      # enum: ["tsr", "perception"]
    output_format: str = "auto", # enum: ["auto", "otsl", "html"]
    max_tokens: int = 8192,
    prompt_text: str = "Extract table from this image."
) -> dict:
    """
    Extract table structure from a table image via VLM backend.

    Returns (JSON):
    {
        "raw_text": str,           # raw model output
        "detected_format": "otsl"|"html",
        "otsl": str,               # present if detected_format=otsl (or output_format allows)
        "html": str,                # present if detected_format=html OR converted from otsl (when
        possible)
        "backend_used": "tsr"|"perception",
        "model": str,
        "base_url": str,
        # optional:
        "error": str
    }
    pass
```

```
def otsl_to_html(otsl_text: str) -> dict:
    """
    Convert OTSL -> HTML.

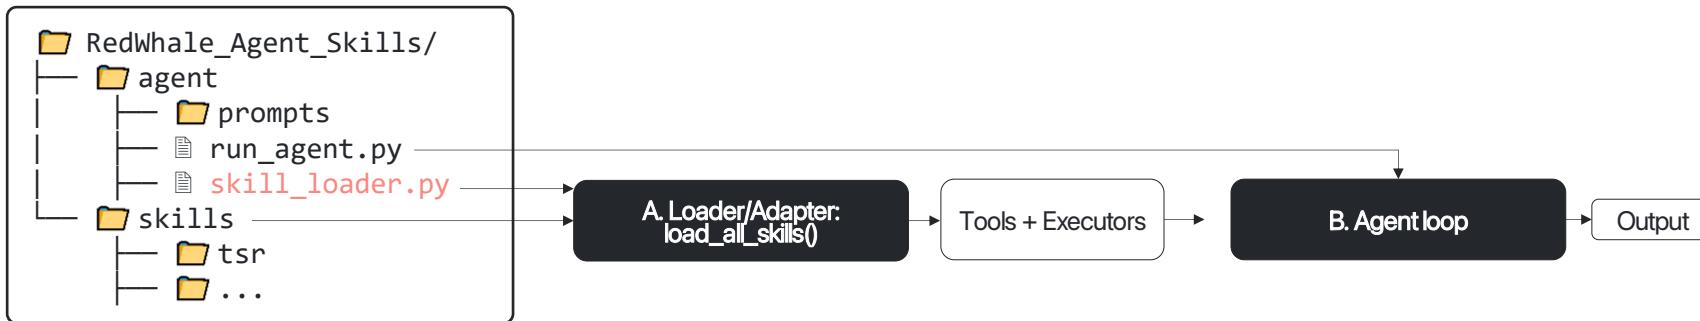
    Returns (JSON):
    {
        "html": str,
        # optional:
        "error": str
    }
    pass
```

```
def save_html(html: str, out_path: str) -> dict:
    """
    Save HTML string to file.

    Returns (JSON):
    {
        "out_path": str,
        # optional:
        "error": str
    }
    pass
```

Implementation

■ How Does It Work? (E2E Flow)



Input: path to skills/
Output:

- tools: a list of tool schemas to be passed into client.responses.create(...)
- runtime["executors"]: a dictionary {tool_name: python_callable} used to actually execute the tools

Mechanism:

- Iterate through each folder inside skills/
- Skip folders that are not valid skills (hidden folders, or those without SKILL.md)
- Parse SKILL.md → extract name and description
- Enter the scripts/directory and import each .py file
- In each script, call the convention function get_registered_tools() to obtain a list of tools:

```
{"name": ..., "description": ..., "parameters": ..., "fn": ...}
```

- Convert them into Responses API schema format:

```
{"type": "function", "name": ..., "description": ..., "parameters": ..., "strict": true}
```

- Store fn into executors for local execution

Result:

- The OpenAI model sees "tools" as a list of functions with schemas
- Your runtime knows which Python function corresponds to each tool

Input: user text
Output: final text response

Loop mechanism (per iteration):

1. Call the model:

```
resp = client.responses.create(model=..., tools=tools, input=input_list)
```
2. Append resp.output to input_list
 - Purpose: to maintain "trace items" (especially for reasoning models)
 - Without this, the model may lose context in the next iteration
3. Filter items with type == "function_call" in resp.output
4. If there is no tool call → end the loop and print resp.output_text
5. For each function_call:
 - Extract name, arguments, call_id
 - Execute the actual Python function:

```
executors[name](**args)
```
 - Append a function_call_output item to return the result to the model:

```
{"type": "function_call_output", "call_id": ..., "output": "...json..."}
```
6. Return to step (1) with the updated input_list containing tool outputs
7. The model reads the tool outputs and synthesizes them into the final answer